

140509_37.md - Privacy-Preserving AI Training Framework

README

Summary: Build a framework that enables AI model training while preserving data privacy through techniques like differential privacy, federated learning, and secure computation.

Problem Statement: Organizations need to train AI models on sensitive data while maintaining privacy and regulatory compliance. Your task is to create a framework that implements privacy-preserving techniques for AI training including differential privacy, federated learning, and secure multi-party computation. The system should provide privacy guarantees, enable collaborative learning without data sharing, and maintain model utility.

Steps: - Design differential privacy mechanisms for various machine learning algorithms - Create secure aggregation and communication protocols for multi-party computation - Build privacy tracking systems - Develop model utility assessment under different privacy constraints - Include compliance reporting and privacy audit capabilities

Suggested Data Requirements: - Sensitive training datasets with privacy requirements - Model performance benchmarks under privacy constraints

Themes: Responsible AI, AI design that assures Security, Legal and Privacy requirements

PRD (Product Requirements Document)

Product Vision

Create a comprehensive privacy-preserving AI training framework that enables organizations to collaborate on AI model development while maintaining strict data privacy and regulatory compliance.

Target Users

- **Primary:** Data Scientists, ML Engineers, Privacy Officers
- **Secondary:** Healthcare Organizations, Financial Institutions, Government Agencies
- **Tertiary:** Research Institutions, Multi-party Collaborations

Core Value Propositions

1. **Privacy Guarantees:** Mathematical privacy guarantees with configurable privacy budgets
2. **Collaborative Learning:** Multi-party AI training without data sharing
3. **Regulatory Compliance:** Built-in compliance with GDPR, HIPAA, CCPA
4. **Utility Preservation:** Minimal impact on model performance while ensuring privacy
5. **Scalable Architecture:** Support for large-scale distributed training

Key Features

1. **Differential Privacy:** Automated DP mechanism design and implementation
2. **Federated Learning:** Secure aggregation and communication protocols
3. **Secure Multi-Party Computation:** Cryptographic protocols for joint computation
4. **Privacy Budget Management:** Automated tracking and optimization
5. **Utility-Privacy Trade-off Analysis:** Comprehensive analysis tools
6. **Compliance Dashboard:** Real-time privacy compliance monitoring

Success Metrics

- Privacy guarantee strength: Configurable ϵ values from 0.1 to 10
 - Model utility preservation: >90% accuracy retention under strong privacy
 - Training efficiency: <2x overhead compared to non-private training
 - Compliance coverage: 100% automated compliance with major privacy laws
 - Adoption rate: 200+ organizations using framework within 12 months
-

FRD (Functional Requirements Document)

Core Functional Requirements

F1: Differential Privacy Implementation

- **F1.1:** Gaussian and Laplacian noise mechanisms for various data types
- **F1.2:** Advanced composition theorems and privacy accounting
- **F1.3:** Private gradient computation with clipping and noise addition
- **F1.4:** Adaptive privacy budget allocation across training epochs
- **F1.5:** Private hyperparameter tuning with privacy budget management

F2: Federated Learning Framework

- **F2.1:** Secure aggregation protocols with cryptographic guarantees
- **F2.2:** Client selection and sampling strategies for heterogeneous data
- **F2.3:** Communication-efficient protocols with compression
- **F2.4:** Byzantine-robust aggregation against malicious participants
- **F2.5:** Personalization techniques for non-IID data distributions

F3: Secure Multi-Party Computation

- **F3.1:** Secret sharing schemes for distributed computation
- **F3.2:** Homomorphic encryption for private arithmetic operations
- **F3.3:** Garbled circuits for complex private functions
- **F3.4:** Private set intersection for data alignment
- **F3.5:** Threshold cryptography for distributed key management

F4: Privacy Budget Management

- **F4.1:** Automated privacy accounting with composition theorems
- **F4.2:** Dynamic privacy budget allocation optimization
- **F4.3:** Privacy budget auditing and compliance reporting
- **F4.4:** Multi-level privacy budgets for hierarchical organizations
- **F4.5:** Privacy budget forecasting and planning tools

F5: Utility-Privacy Analysis

- **F5.1:** Comprehensive utility metrics under privacy constraints
- **F5.2:** Privacy-utility Pareto frontier analysis
- **F5.3:** Sensitivity analysis for privacy parameters
- **F5.4:** Model performance benchmarking across privacy levels
- **F5.5:** Automated privacy parameter optimization

F6: Compliance and Auditing

- **F6.1:** GDPR Article 25 compliance (Privacy by Design)
- **F6.2:** HIPAA Privacy Rule compliance for healthcare data
- **F6.3:** CCPA compliance for consumer data protection
- **F6.4:** Automated privacy audit trail generation
- **F6.5:** Regulatory reporting and documentation automation

NFRD (Non-Functional Requirements Document)

Privacy Requirements

- **NFR-PR1:** Differential privacy guarantees: $\epsilon \in [0.1, 10]$ with $\delta \leq 10^{-6}$
- **NFR-PR2:** Zero-knowledge proofs for computation correctness
- **NFR-PR3:** Information-theoretic security for secret sharing
- **NFR-PR4:** Semantic security for homomorphic encryption
- **NFR-PR5:** Privacy budget consumption tracking with 99.9% accuracy

Performance Requirements

- **NFR-P1:** Training overhead: $<2\times$ compared to non-private training
- **NFR-P2:** Communication efficiency: $<10\text{MB}$ per round in federated setting
- **NFR-P3:** Cryptographic operation latency: $<100\text{ms}$ per operation
- **NFR-P4:** Privacy accounting computation: <1 second per update
- **NFR-P5:** Secure aggregation latency: <30 seconds for 1000 participants

Scalability Requirements

- **NFR-S1:** Support 10,000+ federated learning participants
- **NFR-S2:** Handle datasets up to 1TB with differential privacy
- **NFR-S3:** Multi-party computation with up to 100 parties
- **NFR-S4:** Concurrent training sessions: 100+ simultaneous jobs
- **NFR-S5:** Privacy budget management for 1000+ privacy accounts

Security Requirements

- **NFR-SE1:** End-to-end encryption for all communications
- **NFR-SE2:** Authenticated key exchange protocols
- **NFR-SE3:** Secure random number generation for noise
- **NFR-SE4:** Protection against timing and side-channel attacks
- **NFR-SE5:** Audit logging with integrity guarantees

AD (Architecture Diagram)

```
graph TB
    subgraph "Client Applications"
        PYTHON_SDK[Python SDK]
        R_SDK[R SDK]
        WEB_UI[Web Interface]
        CLI[CLI Tools]
    end

    subgraph "API Gateway & Security"
        GATEWAY[API Gateway]
        AUTH[Authentication]
        AUTHZ[Authorization]
        AUDIT[Audit Logger]
    end

    subgraph "Core Privacy Services"
        DP_ENGINE[Differential Privacy Engine]
        FL_COORDINATOR[Federated Learning Coordinator]
        SMPC_ENGINE[Secure Multi-Party Computation]
        PRIVACY_BUDGET[Privacy Budget Manager]
        UTILITY_ANALYZER[Utility-Privacy Analyzer]
    end

    subgraph "Privacy Mechanisms"
        NOISE_GEN[Noise Generation]
        GRADIENT_CLIP[Gradient Clipping]
        SECURE_AGG[Secure Aggregation]
        SECRET_SHARE[Secret Sharing]
        HE_COMPUTE[Homomorphic Encryption]
    end

    subgraph "Communication Layer"
        SECURE_COMM[Secure Communication]
        KEY_MGMT[Key Management]
        CRYPTO_PROTO[Cryptographic Protocols]
        P2P_NETWORK[P2P Network Layer]
    end

    subgraph "Data Storage"
        ENCRYPTED_DB[Encrypted Database]
        PRIVACY_LOGS[Privacy Audit Logs]
        KEY_STORE[Secure Key Storage]
        MODEL_STORE[Model Repository]
    end

    subgraph "External Integrations"
        COMPLIANCE_SYS[Compliance Systems]
        ML_PLATFORMS[ML Platforms]
        MONITORING[Privacy Monitoring]
        ALERTS[Alert Systems]
    end

    PYTHON_SDK --> GATEWAY
    R_SDK --> GATEWAY
    WEB_UI --> GATEWAY
    CLI --> GATEWAY

    GATEWAY --> AUTH
    GATEWAY --> AUTHZ
    GATEWAY --> AUDIT
```

```
GATEWAY --> DP_ENGINE
GATEWAY --> FL_COORDINATOR
GATEWAY --> SMPC_ENGINE
GATEWAY --> PRIVACY_BUDGET
GATEWAY --> UTILITY_ANALYZER

DP_ENGINE --> NOISE_GEN
DP_ENGINE --> GRADIENT_CLIP
FL_COORDINATOR --> SECURE_AGG
SMPC_ENGINE --> SECRET_SHARE
SMPC_ENGINE --> HE_COMPUTE

FL_COORDINATOR --> SECURE_COMM
SMPC_ENGINE --> KEY_MGMT
SECURE_COMM --> CRYPTO_PROTO
CRYPTO_PROTO --> P2P_NETWORK

DP_ENGINE --> ENCRYPTED_DB
PRIVACY_BUDGET --> PRIVACY_LOGS
KEY_MGMT --> KEY_STORE
UTILITY_ANALYZER --> MODEL_STORE

AUDIT --> COMPLIANCE_SYS
UTILITY_ANALYZER --> ML_PLATFORMS
PRIVACY_BUDGET --> MONITORING
FL_COORDINATOR --> ALERTS
```

HLD (High Level Design)

Differential Privacy Engine

```
class DifferentialPrivacyEngine:
class HomomorphicEncryptionEngine:
    def __init__(self):
        self.key_generator = HEKeyGenerator()
        self.encryptor = HEEncryptor()
        self.evaluator = HEEvaluator()
        self.decryptor = HEDecryptor()

    def private_model_training(self, encrypted_data, model_params, training_config):
        """Train model on homomorphically encrypted data"""

        # Generate homomorphic encryption keys
        public_key, secret_key, evaluation_keys = self.key_generator.generate_keys(
            security_level=training_config.security_level
        )

        # Initialize encrypted model parameters
        encrypted_weights = self.encryptor.encrypt_tensor(
            model_params.weights, public_key
        )

        for epoch in range(training_config.epochs):
            encrypted_gradients = self.compute_encrypted_gradients(
                encrypted_data, encrypted_weights, evaluation_keys
            )

            # Update weights homomorphically
            encrypted_weights = self.evaluator.subtract(
                encrypted_weights,
                self.evaluator.multiply_plain(
                    encrypted_gradients,
                    training_config.learning_rate
                )
            )

        # Decrypt final weights (only by authorized party)
        final_weights = self.decryptor.decrypt_tensor(encrypted_weights, secret_key)

        return HomomorphicTrainingResult(
            trained_weights=final_weights,
            privacy_level='semantic security',
            computational_overhead=self.measure_he_overhead()
        )

    def compute_encrypted_gradients(self, encrypted_data, encrypted_weights, eval_keys):
        """Compute gradients on encrypted data"""
        encrypted_predictions = self.evaluator.matrix_multiply(
            encrypted_data, encrypted_weights, eval_keys
        )
```

```

        # Compute encrypted loss gradients
        encrypted_gradients = self.evaluator.compute_gradient_approximation(
            encrypted_data, encrypted_predictions, eval_keys
        )

        return encrypted_gradients

class SecureAggregationProtocol:
    def __init__(self):
        self.threshold_crypto = ThresholdCryptography()
        self.verifiable_secret_sharing = VerifiableSecretSharing()

    def secure_federated_aggregation(self, client_updates, aggregation_config):
        """Perform secure aggregation of client model updates"""

        num_clients = len(client_updates)
        threshold = aggregation_config.threshold or (num_clients // 2 + 1)

        # Phase 1: Masked model sharing
        client_masks = {}
        shared_updates = {}

        for i, client_update in enumerate(client_updates):
            # Generate random mask
            client_mask = self.generate_random_mask(client_update.shape)
            client_masks[i] = client_mask

            # Mask the client update
            masked_update = client_update + client_mask

            # Secret share the masked update
            shares = self.verifiable_secret_sharing.share_secret(
                masked_update, threshold, num_clients
            )
            shared_updates[i] = shares

        # Phase 2: Mask cancellation
        aggregated_masks = np.zeros_like(client_updates[0])

        for client_id, mask in client_masks.items():
            # Share the mask for cancellation
            mask_shares = self.verifiable_secret_sharing.share_secret(
                mask, threshold, num_clients
            )

            # Reconstruct and subtract mask
            if self.can_reconstruct_secret(mask_shares, threshold):
                reconstructed_mask = self.verifiable_secret_sharing.reconstruct_secret(
                    mask_shares, threshold
                )
                aggregated_masks -= reconstructed_mask

        # Phase 3: Secure aggregation
        aggregated_update = np.zeros_like(client_updates[0])

        for client_id, shares in shared_updates.items():
            if self.can_reconstruct_secret(shares, threshold):
                client_contribution = self.verifiable_secret_sharing.reconstruct_secret(
                    shares, threshold
                )
                aggregated_update += client_contribution

        # Add back the aggregated masks (they should cancel out)
        final_aggregated_update = aggregated_update + aggregated_masks

        return SecureAggregationResult(
            aggregated_update=final_aggregated_update,
            participating_clients=len(client_updates),
            privacy_guarantee='information_theoretic'
        )

class PrivacyBudgetManager:
    def __init__(self):
        self.budget_tracker = BudgetTracker()
        self.composition_accountant = AdvancedCompositionAccountant()

    def allocate_privacy_budget(self, total_budget, training_phases):
        """Optimally allocate privacy budget across training phases"""

        # Analyze the sensitivity of each training phase
        phase_sensitivities = {}

```

```

        for phase in training_phases:
            sensitivity = self.analyze_phase_sensitivity(phase)
            phase_sensitivities[phase.id] = sensitivity

        # Solve optimization problem for budget allocation
        budget_allocation = self.solve_budget_optimization(
            total_budget, phase_sensitivities
        )

        return budget_allocation

def track_privacy_consumption(self, mechanism_type, epsilon, delta):
    """Track privacy budget consumption with advanced composition"""

    # Update composition using Renyi DP or other advanced methods
    self.composition_accountant.compose(
        mechanism=mechanism_type,
        epsilon=epsilon,
        delta=delta
    )

    # Check if budget is exceeded
    current_budget = self.composition_accountant.get_current_budget()

    if current_budget.epsilon > self.budget_tracker.total_epsilon:
        raise PrivacyBudgetExceededException(
            f"Privacy budget exceeded: {current_budget.epsilon} > {self.budget_tracker.total_epsilon}"
        )

    return current_budget

def optimize_noise_parameters(self, target_epsilon, target_delta, utility_constraint):
    """Optimize noise parameters for given privacy and utility constraints"""

    def objective_function(noise_params):
        # Compute expected utility loss
        utility_loss = self.estimate_utility_loss(noise_params)

        # Compute privacy guarantee
        privacy_epsilon = self.compute_privacy_epsilon(noise_params)

        # Penalize if privacy constraint violated
        if privacy_epsilon > target_epsilon:
            return float('inf')

        # Penalize if utility constraint violated
        if utility_loss > utility_constraint:
            return float('inf')

        return utility_loss

    # Use optimization algorithm to find optimal noise parameters
    optimal_params = self.optimize(objective_function)

    return optimal_params

```

Database Schema

```

-- Privacy training jobs
CREATE TABLE privacy_training_jobs (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    job_name VARCHAR(255) NOT NULL,
    privacy_technique VARCHAR(50) NOT NULL, -- 'differential_privacy', 'federated_learning', 'secure_mpc'
    privacy_parameters JSONB NOT NULL,
    dataset_metadata JSONB NOT NULL,
    model_config JSONB NOT NULL,
    status VARCHAR(50) DEFAULT 'created',
    created_by UUID NOT NULL,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    started_at TIMESTAMP,
    completed_at TIMESTAMP,

    CONSTRAINT valid_privacy_technique CHECK (
        privacy_technique IN ('differential_privacy', 'federated_learning', 'secure_mpc', 'homomorphic_encryption')
    )
);

-- Privacy budget tracking
CREATE TABLE privacy_budgets (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    training_job_id UUID REFERENCES privacy_training_jobs(id) ON DELETE CASCADE,

```

```

    total_epsilon DECIMAL(10, 6) NOT NULL,
    total_delta DECIMAL(15, 12) NOT NULL,
    consumed_epsilon DECIMAL(10, 6) DEFAULT 0,
    consumed_delta DECIMAL(15, 12) DEFAULT 0,
    composition_method VARCHAR(50) DEFAULT 'advanced_composition',
    budget_allocation JSONB, -- Per-phase allocation
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    updated_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Privacy mechanism executions
CREATE TABLE privacy_mechanism_executions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    budget_id UUID REFERENCES privacy_budgets(id) ON DELETE CASCADE,
    mechanism_type VARCHAR(100) NOT NULL,
    mechanism_parameters JSONB NOT NULL,
    epsilon_consumed DECIMAL(10, 6) NOT NULL,
    delta_consumed DECIMAL(15, 12) NOT NULL,
    utility_metrics JSONB,
    execution_time_ms INTEGER,
    executed_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

-- Federated learning participants
CREATE TABLE fl_participants (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    training_job_id UUID REFERENCES privacy_training_jobs(id) ON DELETE CASCADE,
    participant_id VARCHAR(255) NOT NULL,
    participant_type VARCHAR(50) NOT NULL, -- 'client', 'server', 'coordinator'
    connection_info JSONB NOT NULL,
    data_characteristics JSONB, -- Size, distribution info
    participation_rounds INTEGER[] DEFAULT '{}',
    status VARCHAR(50) DEFAULT 'registered',
    registered_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    last_seen_at TIMESTAMP
);

-- Secure computation sessions
CREATE TABLE secure_computation_sessions (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    training_job_id UUID REFERENCES privacy_training_jobs(id) ON DELETE CASCADE,
    session_type VARCHAR(50) NOT NULL, -- 'secret_sharing', 'homomorphic_encryption', 'garbled_circuits'
    parties JSONB NOT NULL, -- List of participating parties
    security_parameters JSONB NOT NULL,
    computation_result_hash VARCHAR(64),
    session_status VARCHAR(50) DEFAULT 'initialized',
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    completed_at TIMESTAMP
);

-- Privacy compliance records
CREATE TABLE privacy_compliance_records (
    id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
    training_job_id UUID REFERENCES privacy_training_jobs(id) ON DELETE CASCADE,
    regulation_type VARCHAR(50) NOT NULL, -- 'gdpr', 'hipaa', 'ccpa'
    compliance_status VARCHAR(50) NOT NULL,
    compliance_details JSONB NOT NULL,
    audit_trail JSONB NOT NULL,
    compliance_officer UUID,
    verified_at TIMESTAMP,
    expires_at TIMESTAMP,
    created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);

```

Pseudocode

Differential Privacy Training Workflow

ALGORITHM PrivacyPreservingTraining
 INPUT: model, dataset, privacy_config
 OUTPUT: private_model, privacy_report

```

BEGIN
    // Initialize privacy framework
    privacy_engine = SELECT_PRIVACY_ENGINE(privacy_config.technique)
    privacy_accountant = PrivacyAccountant(privacy_config.epsilon, privacy_config.delta)

    // Validate privacy parameters
    IF NOT VALIDATE_PRIVACY_PARAMETERS(privacy_config) THEN
        RETURN ERROR("Invalid privacy parameters")

```

```

END IF

SWITCH privacy_config.technique
CASE "differential_privacy":
    result = DIFFERENTIAL_PRIVACY_TRAINING(model, dataset, privacy_config, privacy_accountant)
CASE "federated_learning":
    result = FEDERATED_LEARNING_TRAINING(model, dataset, privacy_config, privacy_accountant)
CASE "secure_mpc":
    result = SECURE_MPC_TRAINING(model, dataset, privacy_config, privacy_accountant)
CASE "homomorphic_encryption":
    result = HOMOMORPHIC_ENCRYPTION_TRAINING(model, dataset, privacy_config, privacy_accountant)
DEFAULT:
    RETURN ERROR("Unsupported privacy technique")
END SWITCH

// Generate privacy compliance report
compliance_report = GENERATE_COMPLIANCE_REPORT(result, privacy_config)

// Validate model utility
utility_metrics = EVALUATE_MODEL_UTILITY(result.private_model, dataset)

RETURN PrivacyPreservingResult(
    private_model = result.private_model,
    privacy_spent = privacy_accountant.get_total_privacy_spent(),
    utility_metrics = utility_metrics,
    compliance_report = compliance_report
)
END

FUNCTION DIFFERENTIAL_PRIVACY_TRAINING(model, dataset, config, accountant)
BEGIN
    dp_engine = DifferentialPrivacyEngine()

    // Allocate privacy budget across epochs
    per_epoch_budget = ALLOCATE_PRIVACY_BUDGET(
        total_epsilon = config.epsilon,
        num_epochs = config.num_epochs,
        allocation_strategy = config.budget_allocation
    )

    private_model = model.copy()

    FOR epoch IN RANGE(config.num_epochs) DO
        epoch_epsilon = per_epoch_budget[epoch]

        // Compute gradient sensitivity
        gradient_sensitivity = COMPUTE_GRADIENT_SENSITIVITY(private_model, dataset)

        // Process batches with differential privacy
        epoch_gradients = []

        FOR batch IN dataset.batches(config.batch_size) DO
            // Compute per-sample gradients
            per_sample_gradients = []
            FOR sample IN batch DO
                gradient = private_model.compute_gradient(sample)
                per_sample_gradients.APPEND(gradient)
            END FOR

            // Clip gradients to bound sensitivity
            clipped_gradients = CLIP_GRADIENTS(
                per_sample_gradients,
                max_norm = config.max_gradient_norm
            )

            // Add calibrated Gaussian noise
            noise_scale = COMPUTE_NOISE_SCALE(
                sensitivity = config.max_gradient_norm,
                epsilon = epoch_epsilon / dataset.num_batches,
                delta = config.delta / (config.num_epochs * dataset.num_batches)
            )

            noisy_gradient = ADD_GAUSSIAN_NOISE(
                MEAN(clipped_gradients),
                noise_scale
            )

            epoch_gradients.APPEND(noisy_gradient)

            // Update privacy accountant
            accountant.add_mechanism(
                mechanism = "gaussian_mechanism",

```



```

        epsilon = epoch_epsilon / dataset.num_batches,
        delta = config.delta / (config.num_epochs * dataset.num_batches)
    )
END FOR

// Update model with noisy gradients
aggregated_gradient = MEAN(epoch_gradients)
private_model.update_weights(aggregated_gradient, config.learning_rate)

// Monitor privacy consumption
current_privacy = accountant.get_current_privacy()
IF current_privacy.epsilon > config.epsilon THEN
    BREAK // Stop training if privacy budget exceeded
END IF
END FOR

RETURN DifferentialPrivacyResult(
    private_model = private_model,
    privacy_consumed = accountant.get_current_privacy(),
    training_metrics = GET_TRAINING_METRICS()
)
END

FUNCTION FEDERATED_LEARNING_TRAINING(global_model, client_datasets, config, accountant)
BEGIN
    coordinator = FederatedLearningCoordinator()
    secure_aggregator = SecureAggregator()

    // Initialize federated learning setup
    participants = INITIALIZE_PARTICIPANTS(client_datasets, config)
    global_model_params = global_model.get_parameters()

    FOR round_num IN RANGE(config.num_rounds) DO
        // Select participants for this round
        selected_participants = SELECT_PARTICIPANTS(
            participants,
            selection_fraction = config.client_fraction,
            selection_strategy = config.selection_strategy
        )

        // Distribute current global model to selected participants
        participant_updates = []

        FOR participant IN selected_participants DO
            // Send model to participant
            SEND_MODEL_TO_PARTICIPANT(global_model_params, participant)

            // Participant performs local training with privacy
            local_update = PARTICIPANT_LOCAL_TRAINING(
                participant,
                global_model_params,
                config.local_training_config,
                config.local_privacy_config
            )

            participant_updates.APPEND(local_update)
        END FOR

        // Secure aggregation of participant updates
        IF config.use_secure_aggregation THEN
            aggregated_update = secure_aggregator.secure_aggregate(
                participant_updates,
                aggregation_threshold = config.aggregation_threshold
            )
        ELSE
            aggregated_update = SIMPLE_FEDERATED_AVERAGING(participant_updates)
        END IF

        // Apply differential privacy to aggregated update if configured
        IF config.server_side_dp THEN
            dp_aggregated_update = ADD_DP_NOISE_TO_UPDATE(
                aggregated_update,
                config.server_epsilon / config.num_rounds,
                config.server_delta / config.num_rounds
            )

            accountant.add_mechanism(
                mechanism = "gaussian_mechanism",
                epsilon = config.server_epsilon / config.num_rounds,
                delta = config.server_delta / config.num_rounds
            )
        ELSE

```

```

        dp_aggregated_update = aggregated_update
    END IF

    // Update global model
    global_model_params = UPDATE_GLOBAL_MODEL(
        global_model_params,
        dp_aggregated_update,
        config.server_learning_rate
    )

    // Evaluate global model periodically
    IF round_num % config.evaluation_frequency == 0 THEN
        global_metrics = EVALUATE_GLOBAL_MODEL(
            global_model_params,
            config.validation_data
        )
        LOG_FEDERATED_METRICS(round_num, global_metrics)
    END IF
END FOR

// Finalize global model
final_global_model = CREATE_MODEL_FROM_PARAMETERS(global_model_params)

RETURN FederatedLearningResult(
    global_model = final_global_model,
    total_rounds = config.num_rounds,
    participating_clients = participants.length,
    privacy_consumed = accountant.get_current_privacy(),
    convergence_metrics = GET_CONVERGENCE_METRICS()
)
END

FUNCTION PARTICIPANT_LOCAL_TRAINING(participant, global_params, local_config, privacy_config)
BEGIN
    // Initialize local model with global parameters
    local_model = CREATE_MODEL_FROM_PARAMETERS(global_params)
    local_dataset = participant.get_local_dataset()

    // Apply local differential privacy if configured
    IF privacy_config.use_local_dp THEN
        local_accountant = PrivacyAccountant(
            privacy_config.local_epsilon,
            privacy_config.local_delta
        )

        FOR local_epoch IN RANGE(local_config.local_epochs) DO
            // Compute private gradients
            private_gradient = COMPUTE_PRIVATE_LOCAL_GRADIENT(
                local_model,
                local_dataset,
                privacy_config.local_epsilon / local_config.local_epochs,
                privacy_config.local_delta / local_config.local_epochs
            )

            // Update local model
            local_model.update_weights(private_gradient, local_config.local_learning_rate)

            local_accountant.add_mechanism(
                mechanism = "gaussian_mechanism",
                epsilon = privacy_config.local_epsilon / local_config.local_epochs,
                delta = privacy_config.local_delta / local_config.local_epochs
            )
        END FOR
    ELSE
        // Standard local training without privacy
        FOR local_epoch IN RANGE(local_config.local_epochs) DO
            gradient = local_model.compute_gradient(local_dataset)
            local_model.update_weights(gradient, local_config.local_learning_rate)
        END FOR
    END IF

    // Compute model update (difference from global model)
    model_update = COMPUTE_MODEL_UPDATE(global_params, local_model.get_parameters())

    RETURN ParticipantUpdate(
        update = model_update,
        data_size = local_dataset.size,
        training_loss = local_model.evaluate_loss(local_dataset),
        privacy_consumed = local_accountant.get_current_privacy() IF privacy_config.use_local_dp ELSE None
    )
END

```

```

FUNCTION SECURE_MPC_TRAINING(model, datasets, config, accountant)
BEGIN
    mpc_engine = SecureMultiPartyComputationEngine()

    // Initialize secure computation protocol
    parties = INITIALIZE_MPC_PARTIES(datasets, config)
    protocol = SELECT_MPC_PROTOCOL(config.security_model) // secret_sharing, homomorphic_encryption, garbled_circuits

    // Set up secure communication channels
    secure_channels = ESTABLISH_SECURE_CHANNELS(parties, config.communication_config)

    SWITCH protocol
    CASE "secret_sharing":
        result = SECRET_SHARING_TRAINING(model, parties, config, secure_channels)
    CASE "homomorphic_encryption":
        result = HOMOMORPHIC_ENCRYPTION_TRAINING(model, parties, config, secure_channels)
    CASE "garbled_circuits":
        result = GARBLED_CIRCUITS_TRAINING(model, parties, config, secure_channels)
    END SWITCH

    RETURN SecureMPCResult(
        trained_model = result.model,
        computation_transcript = result.transcript,
        security_guarantees = result.security_level,
        performance_metrics = result.performance
    )
END

FUNCTION SECRET_SHARING_TRAINING(model, parties, config, channels)
BEGIN
    secret_sharing_scheme = SELECT_SECRET_SHARING_SCHEME(config.threshold, parties.length)

    // Phase 1: Share training data
    shared_datasets = {}
    FOR party IN parties DO
        data_shares = secret_sharing_scheme.share_data(
            party.dataset,
            threshold = config.threshold,
            num_parties = parties.length
        )
        shared_datasets[party.id] = data_shares
    END FOR

    // Phase 2: Initialize shared model parameters
    shared_model_params = secret_sharing_scheme.share_data(
        model.get_parameters(),
        threshold = config.threshold,
        num_parties = parties.length
    )

    // Phase 3: Secure training loop
    FOR epoch IN RANGE(config.num_epochs) DO
        // Compute gradients on shared data
        shared_gradients = COMPUTE_GRADIENTS_ON_SHARES(
            shared_model_params,
            shared_datasets,
            secret_sharing_scheme
        )

        // Update shared model parameters
        shared_model_params = UPDATE_SHARED_PARAMETERS(
            shared_model_params,
            shared_gradients,
            config.learning_rate,
            secret_sharing_scheme
        )
    END FOR

    // Phase 4: Reconstruct final model
    IF PARTIES_AGREE_TO_RECONSTRUCT(parties, config.threshold) THEN
        final_model_params = secret_sharing_scheme.reconstruct_secret(
            shared_model_params,
            threshold = config.threshold
        )
        final_model = CREATE_MODEL_FROM_PARAMETERS(final_model_params)
    ELSE
        final_model = NULL // Cannot reconstruct without sufficient parties
    END IF

    RETURN SecretSharingResult(
        model = final_model,
        privacy_level = "information_theoretic",

```

```

        reconstruction_successful = (final_model IS NOT NULL)
    )
END

```

This completes the comprehensive documentation for Problem Statements 36 and 37. Both solutions provide enterprise-grade architectures for responsible AI, covering bias detection/mitigation and privacy-preserving training frameworks with advanced cryptographic protocols and regulatory compliance.): self.noise_generator = NoiseGenerator() self.privacy_accountant = PrivacyAccountant() self.gradient_clipper = GradientClipper() self.composition_tracker = CompositionTracker()

```

def private_training(self, model, dataset, privacy_config):
    # Initialize privacy parameters
    epsilon = privacy_config.epsilon
    delta = privacy_config.delta
    epochs = privacy_config.epochs

    # Allocate privacy budget across epochs
    per_epoch_epsilon = self.privacy_accountant.allocate_budget(
        total_epsilon=epsilon,
        total_epochs=epochs,
        allocation_strategy=privacy_config.allocation_strategy
    )

    private_model = model.copy()

    for epoch in range(epochs):
        # Compute private gradients
        private_gradients = self.compute_private_gradients(
            private_model,
            dataset,
            per_epoch_epsilon[epoch],
            delta / epochs
        )

        # Update model with private gradients
        private_model.update(private_gradients)

        # Track privacy consumption
        self.composition_tracker.add_mechanism(
            epsilon=per_epoch_epsilon[epoch],
            delta=delta / epochs
        )

        # Monitor utility degradation
        utility_loss = self.evaluate_utility_loss(private_model, model, dataset)

    return PrivateTrainingResult(
        model=private_model,
        privacy_spent=self.composition_tracker.get_total_privacy(),
        utility_metrics=self.compute_utility_metrics(private_model, dataset)
    )

def compute_private_gradients(self, model, dataset, epsilon, delta):
    batch_size = len(dataset) // 10 # Reasonable batch size
    sensitivity = self.compute_gradient_sensitivity(model, dataset)

    private_gradients = []

    for batch in dataset.batch(batch_size):
        # Compute per-sample gradients
        per_sample_grads = model.compute_per_sample_gradients(batch)

        # Clip gradients to bound sensitivity
        clipped_grads = self.gradient_clipper.clip_gradients(
            per_sample_grads,
            clip_norm=sensitivity
        )

        # Add calibrated noise
        noise_scale = self.compute_noise_scale(sensitivity, epsilon, delta)
        noisy_grads = self.noise_generator.add_gaussian_noise(
            clipped_grads,
            noise_scale
        )

        private_gradients.append(noisy_grads)

    return np.mean(private_gradients, axis=0)

```

```

class FederatedLearningCoordinator:
    def __init__(self):
        self.secure_aggregator = SecureAggregator()
        self.client_manager = ClientManager()
        self.communication_protocol = SecureCommunicationProtocol()

```

```

def coordinate_federated_training(self, global_model, clients, fl_config):
    for round_num in range(fl_config.num_rounds):
        # Select clients for this round
        selected_clients = self.client_manager.select_clients(
            clients,
            selection_strategy=fl_config.client_selection,
            fraction=fl_config.client_fraction
        )

        # Distribute global model to selected clients
        client_updates = []

        for client in selected_clients:
            # Send model to client
            client_model = self.communication_protocol.send_model(
                global_model, client
            )

            # Client performs local training
            local_update = client.local_training(
                client_model,
                fl_config.local_epochs,
                fl_config.local_privacy_config
            )

            client_updates.append(local_update)

        # Secure aggregation of client updates
        aggregated_update = self.secure_aggregator.aggregate_updates(
            client_updates,
            aggregation_weights=self.compute_aggregation_weights(selected_clients)
        )

        # Update global model
        global_model.apply_update(aggregated_update)

        # Evaluate global model
        if round_num % fl_config.eval_frequency == 0:
            global_metrics = self.evaluate_global_model(global_model, fl_config.test_data)

    return FederatedLearningResult(
        global_model=global_model,
        training_metrics=self.get_training_history(),
        privacy_metrics=self.compute_privacy_metrics()
    )

```

```

class SecureMultiPartyComputation:
    def init(self):
        self.secret_sharing = SecretSharingScheme()
        self.homomorphic_encryption = HomomorphicEncryption()
        self.garbled_circuits = GarbledCircuits()

```

```

def secure_joint_training(self, parties, training_function, security_config):
    # Initialize secure computation protocol
    if security_config.protocol == 'secret_sharing':
        return self.secret_sharing_protocol(parties, training_function)
    elif security_config.protocol == 'homomorphic_encryption':
        return self.homomorphic_encryption_protocol(parties, training_function)
    elif security_config.protocol == 'garbled_circuits':
        return self.garbled_circuits_protocol(parties, training_function)

```

```

def secret_sharing_protocol(self, parties, training_function):
    # Share data using secret sharing
    shared_data = []
    for party in parties:
        shares = self.secret_sharing.share_data(
            party.data,
            threshold=len(parties)//2 + 1,
            num_parties=len(parties)
        )
        shared_data.append(shares)

    # Perform secure computation on shared data
    result_shares = training_function.compute_on_shares(shared_data)

    # Reconstruct final result
    final_result = self.secret_sharing.reconstruct(
        result_shares,
        threshold=len(parties)//2 + 1
    )

    return SecureComputationResult(
        result=final_result,
        privacy_guarantees='information_theoretic',
        computation_overhead=self.measure_overhead()
    )

```

```
)  
...  

```

LLD (Low Level Design)

Advanced Differential Privacy Mechanisms

```
``python class AdvancedDPMechanisms: def init(self): self.composition_accountant = RenyiDPAccountant()  
self.sensitivity_analyzer = SensitivityAnalyzer()  
  
def private_sgd(self, model, dataset, privacy_params):  
    """Private Stochastic Gradient Descent with advanced composition"""  
    epsilon, delta = privacy_params.epsilon, privacy_params.delta  
  
    # Compute sensitivity  
    l2_sensitivity = self.sensitivity_analyzer.compute_l2_sensitivity(model)  
  
    # Initialize privacy accountant  
    self.composition_accountant.initialize(epsilon, delta)  
  
    for epoch in range(privacy_params.epochs):  
        for batch in dataset.batches(privacy_params.batch_size):  
            # Compute per-example gradients  
            per_example_grads = []  
            for example in batch:  
                grad = model.compute_gradient(example)  
                per_example_grads.append(grad)  
  
            # Clip gradients  
            clipped_grads = self.clip_gradients_per_example(  
                per_example_grads,  
                privacy_params.max_grad_norm  
            )  
  
            # Compute noise scale using current privacy budget  
            current_epsilon = self.composition_accountant.get_current_epsilon()  
            noise_multiplier = self.compute_noise_multiplier(  
                target_epsilon=current_epsilon,  
                target_delta=delta,  
                sensitivity=privacy_params.max_grad_norm  
            )  
  
            # Add Gaussian noise  
            noisy_grad = self.add_gaussian_noise(  
                np.mean(clipped_grads, axis=0),  
                noise_multiplier * privacy_params.max_grad_norm  
            )  
  
            # Update model  
            model.update_weights(noisy_grad)  
  
            # Update privacy accountant  
            self.composition_accountant.step(noise_multiplier)  
  
        return model, self.composition_accountant.get_privacy_spent()  
  
def private_query_release(self, dataset, queries, privacy_budget):  
    """Private query release using exponential mechanism"""  
    query_results = []  
    epsilon_per_query = privacy_budget / len(queries)  
  
    for query in queries:  
        # Compute utility function  
        def utility_function(answer):  
            return -abs(query.true_answer(dataset) - answer)  
  
        # Sample from exponential mechanism  
        private_answer = self.exponential_mechanism(  
            dataset,  
            query.answer_domain,  
            utility_function,  
            epsilon_per_query / 2 # Global sensitivity of utility  
        )  
  
        query_results.append(private_answer)  
  
    return query_results
```

```
class HomomorphicEncryptionEngine: def init(self
```