

# 140509\_46.md – AI-Enhanced Code Generation and Review Platform

**Theme:** AI in Software Engineering Lifecycle

**Mission:** Boost developer productivity and code quality via context-aware NL→Code generation, automated tests, intelligent code review, and rich IDE/CI integrations.

---

## README (Problem Statement)

**Summary:** Build a comprehensive platform that assists developers with code generation from natural language, automated testing, and intelligent code review across the development lifecycle.

**Problem Statement:** Deliver a platform that understands project context, integrates with IDEs and CI/CD, and suggests code aligned with standards while generating tests and performing high-signal reviews.

**Steps:**

- Natural language to code with context awareness
- Automated test generation & coverage analysis
- Intelligent review (bug/risk detection, optimizations)
- IDE plugins & workflow integrations
- Documentation generation/maintenance
- Quality metrics & technical debt assessment

**Suggested Data:** Large multi-language repos with tests and review comments; style guides; bug/patch histories.

---

## 1) Vision, Scope, KPIs

**Vision:** Make high-quality software the default by embedding AI across design-build-test-review stages.

**Scope:**

- v1: NL→Code, IDE plugin, static checks, unit test stubs, CI comments.
- v2: ML bug detector, integration tests, doc generation, refactoring suggestions.
- v3: Multi-repo context, architectural reviews, technical debt analytics.

**KPIs:**

- Suggestion acceptance rate ≥ 50%
  - Auto-tests raise coverage ≥ 70% lines/branches on new code
  - Review engine catches ≥ 80% of seeded bug patterns
  - Dev cycle time ↓ 30% for target teams
- 

## 2) Personas & User Stories

- **Developer:** Inline NL prompts → code; quick fix & refactor suggestions.
- **QA Engineer:** Auto-generated tests with reports & coverage gates.
- **Tech Lead:** Review dashboards, policy gates, debt trendlines.
- **Security Engineer:** SAST/secret scans with autofixes and PR annotations.

**Stories:**

- US-01: Generate a typed API client from an OpenAPI spec.
  - US-06: Propose tests to cover edge cases identified by symbolic execution.
  - US-12: PR review auto-flags SQL injection and suggests a parameterized fix.
- 

## 3) PRD (Capabilities)

1. **Context-Aware NL→Code:**

- Retrieves relevant files/snippets, types, API usage, and project style; supports

Python/TS/Go/Java/C# in v1.

2. **Test Generation & Coverage:**
    - Unit/integration test synthesis (property-based where applicable); coverage & mutation testing reports.
  3. **Intelligent Review:**
    - Static + ML: bug risk scores, concurrency/safety checks, performance tips, security rules (CWE/OWASP).
  4. **IDE/CI Integration:**
    - VS Code/JetBrains plugins; inline diffs; CI bots with PR annotations and auto-fix PRs.
  5. **Documentation:**
    - Docstrings, READMEs, architecture digests, change logs.
  6. **Quality Metrics:**
    - Lint/complexity/duplication, coverage, hotspot detection, debt scoring.
- 


## 4) FRD (Functional Requirements)

- **RAG Context Server:** AST & symbol index, vector index over code/comments, dependency graph.
  - **Prompt Builder:** Templates inject context (API signatures, tests failing, style rules).
  - **Generators:** CodeGen, TestGen, DocGen with safety rails (no destructive shell, no secrets).
  - **Review Pipeline:** SAST (Bandit/ESLint/semgrep), license checks, secret detection, ML classifiers for bug patterns.
  - **CI Gates:** enforce min coverage, max complexity, security severities; staged approvals.
  - **Policy Engine:** org/team rules (e.g., forbid eval/exec).
  - **Telemetry:** acceptance, regressions, IDE latency; opt-in privacy.
- 

## 5) NFRD

- **Latency:** P95 suggestion  $\approx 300$  ms (cached context);  $\approx 1.5$  s cold.
  - **Scale:** Repos up to 10M LOC; multi-repo context.
  - **Security:** On-prem isolation; no code leaves tenant; SBOM for components.
  - **Reliability:** 99.9% plugin service uptime.
  - **Compliance:** SOC2/ISO27001; code retention policies.
- 

## 6) Architecture (Logical)



## 7) HLD (Key Components)

- **Context Server:**
    - Build AST, symbol table, call graph; compute embeddings per symbol/file; elastic code search.
  - **CodeGen:**
    - Large code LLM; decoding constrained by types & lints; temperature  $\approx 0.2$  by default.
  - **TestGen:**
    - Path exploration (symbolic execution) + heuristics; property-based tests for pure functions.
  - **Review Engine:**
    - Semgrep rules + ML risk model; taint analysis for sinks (SQL, SSRF, command).
  - **DocGen:**
    - Generate docstrings from AST; summarize modules; Mermaid UML/sequence diagrams.
  - **CI Bot:**
    - PR annotations, auto-fix patch generation, rollback/patch explainers.
-

## 8) LLD (Selected)

**Context Retrieval:** - Build query with current file, cursor scope, imported types; fetch top-k symbols from vector index; include failing tests and lint findings.

### Prompt Template (Python):

System: You are a senior Python engineer.  
Context: <snippets+APIs+style+tests>  
Task: Implement function {name} satisfying docstring and tests.  
Constraints: PEP8, type hints, no external calls, raise ValueError on invalid input.

### Review Rule (Semgrep):

```
rules:  
- id: py.sql.injection.param  
  pattern: cursor.execute($QUERY)  
  message: Use parameterized queries.  
  severity: ERROR
```

**CI Gate (Coverage):** - Fail PR if new/changed lines coverage < 70%.

---

## 9) Pseudocode (End-to-End)

```
on_ide_request(prompt, cursor):  
    ctx = retrieve_context(repo, cursor)  
    code = codegen(prompt, ctx)  
    tests = testgen(code, ctx)  
    review = review_engine(code, ctx)  
    docs = docgen(code, ctx)  
    return bundle(code, tests, review, docs)  
  
on_ci_pull_request(pr):  
    metrics = run_checks(pr)  
    if metrics.coverage < 0.7 or metrics.security.high > 0:  
        annotate(pr, metrics)  
        if can_autofix(metrics): create_autofix_pr(pr)  
    else:  
        approve(pr)
```

---

## 10) Data & Evaluation

- **Training/Seeds:** BigCode/The Stack (filtered), CodeSearchNet, internal corpora with consent; review datasets (MSR, Google, GitHub PRs).
  - **Metrics:** suggestion acceptance, edit distance to final, test coverage uplift, bug detection precision/recall, time-to-merge.
  - **A/B:** team-level rollouts; guarded promotion via gates.
- 

## 11) Security & Governance

- PII/secret scrubbing; local inference option; reproducible builds; signed models; audit logs.
  - License compliance checks; third-party component SBOMs.
- 

## 12) Observability & Cost

- Metrics: IDE latency, acceptance %, test gen time, CI queue times, GPU utilization.
  - Cost controls: distillation, quantization, shared KV cache, batching; autoscale.
- 

## 13) Roadmap

- **M1 (4w):** IDE plugin + NLâ†’Code + static checks.

- **M2 (8w):** TestGen + CI bot + coverage gates.
  - **M3 (12w):** ML bug model + DocGen + auto-fixes.
  - **M4 (16w):** Multi-language scale + architectural reviews + debt analytics.
- 

## 14) Risks & Mitigations

- **Hallucinated code:** retrieval augmentation, constrained decoding, unit-test-first mode.
- **False positives in review:** precision-tuned rules, allow suppressions, human-in-loop.
- **Latency spikes:** warm pools, KV cache, local models.
- **IP concerns:** on-prem sealed deployment, data minimization.